

# Big Data

Student: Teodora Stamenković 1460

Profesor: Dragan Stojanović

---

Putanje za 442 taksi vozila u Portu u periodu od 01.07.2013. do 30.06.2014.

- TRIP\_ID
- CALL\_TYPE:
  - A - poziv iz centrale
  - B - direktno kontaktiranje taksija na stajalištu
  - C - ostalo
- ORIGIN\_CALL - identifikacija putnika ukoliko je CALL\_TYPE='A'
- ORIGIN\_STAND - identifikacija taksi stajališta ukoliko je CALL\_TYPE='B'
- TAXI\_ID - identifikacija taksi vozača
- TIMESTAMP - vreme početka
- DAY\_TYPE
  - B - praznik
  - C - dan pre praznika
  - A - obican dan
- MISSING\_DATA - da li nedostaju GPS podaci
- POLYLINE - lista parova (lat, lon) za svakih 15s putovanja

```
root
|-- TRIP_ID: long (nullable = true)
|-- CALL_TYPE: string (nullable = true)
|-- ORIGIN_CALL: integer (nullable = true)
|-- ORIGIN_STAND: integer (nullable = true)
|-- TAXI_ID: integer (nullable = true)
|-- TIMESTAMP: integer (nullable = true)
|-- DAY_TYPE: string (nullable = true)
|-- MISSING_DATA: boolean (nullable = true)
|-- POLYLINE: string (nullable = true)
```

# Prikaz podataka

TRIP_ID	CALL_TYPE	ORIGIN_CALL	ORIGIN_STAND	TAXI_ID	TIMESTAMP	DAY_TYPE	MISSING_DATA	POLYLINE
1372636858620000589	C	null	null	20000589	1372636858	A	false	[[-8.618643,41.14...
1372637303620000596	B	null	7	20000596	1372637303	A	false	[[-8.639847,41.15...
1372636951620000320	C	null	null	20000320	1372636951	A	false	[[-8.612964,41.14...
1372636854620000520	C	null	null	20000520	1372636854	A	false	[[-8.574678,41.15...
1372637091620000337	C	null	null	20000337	1372637091	A	false	[[-8.645994,41.18...
1372636965620000231	C	null	null	20000231	1372636965	A	false	[[-8.615502,41.14...
1372637210620000456	C	null	null	20000456	1372637210	A	false	[[-8.57952,41.145...
1372637299620000011	C	null	null	20000011	1372637299	A	false	[[-8.617563,41.14...
1372637274620000403	C	null	null	20000403	1372637274	A	false	[[-8.611794,41.14...
1372637905620000320	C	null	null	20000320	1372637905	A	false	[[-8.615907,41.14...

only showing top 10 rows

DataFrame Rows count : 1710670

# Transformacija podataka

```
#drop missing values
dataset = dataset.filter(dataset["missing_data"] == False)
dataset = dataset.withColumnRenamed("timestamp", "start_time")

# check missing values
print(dataset.where(dataset["start_time"].isNull()).count())

# data transformation
dataset = dataset.withColumn("coordinates", F.regexp_replace("polyline", "[\\|\\|]", ""))

dataset = dataset.withColumn("start_lon", F.split(dataset["coordinates"], ",").getItem(0).cast('double'))
dataset = dataset.withColumn("start_lat", F.split(dataset["coordinates"], ",").getItem(1).cast('double'))
dataset = dataset.withColumn("end_lon", F.reverse(F.split(dataset["coordinates"], ",").getItem(1).cast('double'))
dataset = dataset.withColumn("end_lat", F.reverse(F.split(dataset["coordinates"], ",").getItem(0).cast('double'))

dataset = dataset.withColumn("array_of_coordinates", F.split(dataset["coordinates"], ","))
dataset = dataset.withColumn("trip_duration", F.size(F.col("array_of_coordinates")) * 7.5 )

dataset = dataset.withColumn("end_time", dataset["start_time"] + dataset["trip_duration"])
dataset = dataset.drop("polyline")
dataset = dataset.drop("coordinates")
dataset = dataset.drop("array_of_coordinates")

dataset.show(10)
```

# Prikaz podataka nakon transformacija

trip_id	call_type	origin_call	origin_stand	taxi_id	start_time	day_type	missing_data	start_lon	start_lat	end_lon	end_lat	trip_duration	end_time
1372636858620000589	C	null	null	20000589	1372636858	A	false	-8.618643	41.141412	-8.630838	41.154489	345.0	1.372637203E9
1372637303620000596	B	null	7	20000596	1372637303	A	false	-8.639847	41.159826	-8.66574	41.170671	285.0	1.372637588E9
1372636951620000320	C	null	null	20000320	1372636951	A	false	-8.612964	41.140359	-8.61597	41.14053	975.0	1.372637926E9
1372636854620000520	C	null	null	20000520	1372636854	A	false	-8.574678	41.151951	-8.607996	41.142915	645.0	1.372637499E9
1372637091620000337	C	null	null	20000337	1372637091	A	false	-8.645994	41.18049	-8.687268	41.178087	435.0	1.372637526E9
1372636965620000231	C	null	null	20000231	1372636965	A	false	-8.615502	41.140674	-8.578224	41.160717	390.0	1.372637355E9
1372637210620000456	C	null	null	20000456	1372637210	A	false	-8.57952	41.145948	-8.603973	41.142816	540.0	1.372637775E9
1372637299620000011	C	null	null	20000011	1372637299	A	false	-8.617563	41.146182	-8.6247	41.161554	510.0	1.372637809E9
1372637274620000403	C	null	null	20000403	1372637274	A	false	-8.611794	41.140557	-8.589402	41.163309	570.0	1.372637844E9
1372637905620000320	C	null	null	20000320	1372637905	A	false	-8.615907	41.140557	-8.604594	41.134158	285.0	1.37263819E9

only showing top 10 rows

# Projekat 1

- Apache Spark aplikacija
  - HDFS
-



# Filtriranje podataka

```
first_latitude = float(os.getenv('FIRST_LATITUDE'))
first_longitude = float(os.getenv('FIRST_LONGITUDE'))
second_latitude = float(os.getenv('SECOND_LATITUDE'))
second_longitude = float(os.getenv('SECOND_LONGITUDE'))
start_time = int(os.getenv('START_TIME'))
end_time = int(os.getenv('END_TIME'))
```

Podaci se filtriraju na osnovu ulaznih podataka (opseg koordinata i vremenski period)

```
dataset_filtered = dataset.filter(((dataset["start_lat"] > first_latitude) & (dataset["start_lon"] > first_longitude)
& (dataset["start_lat"] < second_latitude) & (dataset["start_lon"] < second_longitude)
& (dataset["start_time"] > start_time) & (dataset["start_time"] < end_time))
| ((dataset["end_lat"] > first_latitude) & (dataset["end_lon"] > first_longitude)
& (dataset["end_lat"] < second_latitude) & (dataset["end_lon"] < second_longitude)
& (dataset["end_time"] > start_time) & (dataset["end_time"] < end_time)))
```

# Analiza podataka

```
print("Average trip duration grouped by call type")
dataset_trip_duration_by_call_type = dataset_filtered.groupBy("call_type").agg(F.avg("trip_duration"))
dataset_trip_duration_by_call_type.show()
```

Average trip duration grouped by call type

call_type	avg(trip_duration)
B	838.3928571428571
C	1038.75
A	843.75



# Analiza podataka

```
trip_duration_stddev = dataset_filtered.groupBy("origin_stand").agg(F.stddev("trip_duration").alias("trip_duration_stddev"))
trip_duration_stddev = trip_duration_stddev.sort("trip_duration_stddev", ascending=False)
origin_stand = trip_duration_stddev.collect()[0]
result = "Taxi stand with the widest range of trip duration is " + str(origin_stand.asDict()["origin_stand"]) \
+ " (stddev = " + str(origin_stand.asDict()['trip_duration_stddev']) + ")\n"
```

Output:

```
1 Taxi stand with the widest range of trip duration is 59 (stddev = 1245.2054778031106)
```

# Analiza podataka

```
dataset_max = dataset_filtered.filter(dataset_filtered["call_type"] == 'B').groupBy('origin_stand').agg(count('trip_id')\
    .alias('num_of_trips'), max('trip_duration').alias('max_trip_time'))
dataset_max_duration = dataset_max.sort("max_trip_time", ascending=False)
dataset_max_count = dataset_max.sort("num_of_trips", ascending=False)
max_duration = dataset_max_duration.collect()[0]
max_count = dataset_max_count.collect()[0]
result = "The longest trip was " + str(max_duration.asDict()['max_trip_time'] / 3600) + " hours from taxi stand " \
    + str(max_duration.asDict()['origin_stand']) + "\n"
result_count = "The highest number of trips (" + str(max_count.asDict()['num_of_trips']) + ") started from taxi stand " \
    + str(max_count.asDict()['origin_stand']) + "\n"
```

Output:

```
1 The longest trip was 0.9041666666666667 hours from taxi stand 15
```

```
3 The highest number of trips (19) started from taxi stand 15
```

# Analiza podataka

```
dataset_taxi = dataset_filtered.groupBy("taxi_id").agg(F.sum("trip_duration").alias('trip_duration_sum'))
dataset_taxi = dataset_taxi.filter(dataset_taxi["trip_duration_sum"] > 0).sort("trip_duration_sum", ascending=True)
taxi_driver = dataset_taxi.collect()[0]
result = 'Taxi driver who spent the least time driving is ' + str(taxi_driver.asDict()['taxi_id']) \
         + ' (' + str(taxi_driver.asDict()['trip_duration_sum'] / 60) + ' minutes)\n'
```

Output:

```
4 Taxi driver who spent the least time driving is 20000472 (0.75 minutes)
```

# Čuvanje podataka na HDFS




localhost:9870/explorer.html#/dir

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

## Browse Directory


/dir

Go!



Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	-rw-r--r--	root	supergroup	1.81 GB	Jan 26 23:13	1	128 MB	<a href="#">train.csv</a> 

Showing 1 to 1 of 1 entries

Previous

1

Next

Hadoop, 2019.

# Lokalno izvršenje

```
input = "hdfs://localhost:9000/dir/train.csv"
spark = SparkSession.builder.appName(appName).master("local[2]").getOrCreate()
```

```
dataset = spark.read.option("inferSchema", True).option("header", True).csv(input)
```

## Spark Jobs (?)

User: root  
Total Uptime: 1.8 min  
Scheduling Mode: FIFO  
Active Jobs: 1  
Completed Jobs: 9

Event Timeline  
☐ Enable zooming



# Lokalno izvršenje - executors


← → ↻

localhost:4040/executors/

☆

🔖

☰

 3.3.1

Jobs

Stages

Storage

Environment

Executors

SQL / DataFrame

Taxi Porto application UI

## Executors

[Show Additional Metrics](#)

### Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	88.4 KiB / 434.4 MiB	0.0 B	2	3	0	41	44	35 s (0.4 s)	4.6 GiB	885 B	1.3 KiB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	88.4 KiB / 434.4 MiB	0.0 B	2	3	0	41	44	35 s (0.4 s)	4.6 GiB	885 B	1.3 KiB	0

### Executors

Show  entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.0.12:42305	Active	0	88.4 KiB / 434.4 MiB	0.0 B	2	3	0	41	44	35 s (0.4 s)	4.6 GiB	885 B	1.3 KiB	<a href="#">Thread Dump</a>

Showing 1 to 1 of 1 entries

Previous **1** Next



# Lokalno izvršenje - stages

## Completed Stages (13)

Page: 1

1 Pages. Jump to 1. Show 100 items in a page. Go

Stage Id ▼	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
18	<a href="#">collect at /home/tea_/_bigdata-project/BigData/app/app.py:85</a>	<a href="#">+details</a>	2023/01/30 00:09:09	80 ms	1/1			3.9 KiB	
15	<a href="#">collect at /home/tea_/_bigdata-project/BigData/app/app.py:85</a>	<a href="#">+details</a>	2023/01/30 00:09:09	56 ms	1/1			64.9 KiB	3.9 KiB
13	<a href="#">collect at /home/tea_/_bigdata-project/BigData/app/app.py:85</a>	<a href="#">+details</a>	2023/01/30 00:09:08	68 ms	1/1			64.9 KiB	
11	<a href="#">collect at /home/tea_/_bigdata-project/BigData/app/app.py:85</a>	<a href="#">+details</a>	2023/01/30 00:08:18	50 s	15/15	1853.7 MiB			64.9 KiB
10	<a href="#">showString at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:08:18	20 ms	1/1			3.4 KiB	
8	<a href="#">showString at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:07:26	52 s	15/15	1853.7 MiB			3.4 KiB
7	<a href="#">showString at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:07:25	0.7 s	1/1	2.4 MiB			
6	<a href="#">showString at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:07:25	29 ms	1/1	64.0 KiB			
5	<a href="#">count at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:07:24	60 ms	1/1			840.0 B	
3	<a href="#">count at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:07:14	10 s	15/15	1853.7 MiB			840.0 B
2	<a href="#">showString at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:07:13	79 ms	1/1	64.0 KiB			
1	<a href="#">csv at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:06:59	14 s	15/15	1853.7 MiB			
0	<a href="#">csv at NativeMethodAccessorImpl.java:0</a>	<a href="#">+details</a>	2023/01/30 00:06:58	0.3 s	1/1	64.0 KiB			

# Izvršenje na klasteru

```
input = "hdfs://namenode:9000/dir/train.csv"
spark = SparkSession.builder.appName(appName).master("spark://spark-master:7077").getOrCreate()
```

```
dataset = spark.read.option("inferSchema", True).option("header", True).csv(input)
```



3.1.2

Spark Master at spark://b115afcb22a3:7077

URL: spark://b115afcb22a3:7077

Alive Workers: 2

Cores in use: 8 Total, 8 Used

Memory in use: 13.5 GiB Total, 2.0 GiB Used

Resources in use:

Applications: 1 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

## Workers (2)

Worker Id	Address	State	Cores	Memory	Resources
<a href="#">worker-20230130072909-172.18.0.8-40165</a>	172.18.0.8:40165	ALIVE	4 (4 Used)	6.8 GiB (1024.0 MiB Used)	
<a href="#">worker-20230130072909-172.18.0.9-39879</a>	172.18.0.9:39879	ALIVE	4 (4 Used)	6.8 GiB (1024.0 MiB Used)	

## Running Applications (1)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
<a href="#">app-20230130072940-0000</a>	(kill) <a href="#">Taxi Porto</a>	8	1024.0 MiB		2023/01/30 07:29:40	root	RUNNING	4.9 min

# Izvršenje na klasteru



3.1.2

Jobs

Stages

Storage

Environment

Executors

SQL

Taxi Porto application UI

## Executors

[Show Additional Metrics](#)

### Summary

	▲ RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(3)	0	547.3 KiB / 1.1 GiB	0.0 B	8	6	0	1519	1525	46 min (2.7 min)	11.9 GiB	4.8 KiB	5.5 KiB	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(3)	0	547.3 KiB / 1.1 GiB	0.0 B	8	6	0	1519	1525	46 min (2.7 min)	11.9 GiB	4.8 KiB	5.5 KiB	0

### Executors

Show 20 entries

Search:

Executor ID ▲	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Logs	Thread Dump
0	172.18.0.9:42115	Active	0	164.2 KiB / 366.3 MiB	0.0 B	4	4	0	763	767	24 min (1.5 min)	6.4 GiB	1.4 KiB	1.5 KiB	<a href="#">stdout</a> <a href="#">stderr</a>	<a href="#">Thread Dump</a>
driver	0e07af938c64:35155	Active	0	191.6 KiB / 366.3 MiB	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B		<a href="#">Thread Dump</a>
1	172.18.0.8:39993	Active	0	191.6 KiB / 366.3 MiB	0.0 B	4	2	0	756	758	22 min (1.2 min)	5.5 GiB	3.4 KiB	4 KiB	<a href="#">stdout</a> <a href="#">stderr</a>	<a href="#">Thread Dump</a>

Showing 1 to 3 of 3 entries

Previous 1 Next

# Projekt 2

- Kafka producer
  - Apache Spark
  - Apache Flink
-

# Kafka

```
zookeeper:
  image: 'bitnami/zookeeper:latest'
  container_name: zookeeper
  ports:
    - '2181:2181'
  environment:
    - ALLOW_ANONYMOUS_LOGIN=yes

kafka:
  image: 'bitnami/kafka:latest'
  container_name: kafka
  ports:
    - '9092:9092'
    - '29092:29092'
  environment:
    - KAFKA_BROKER_ID=1
    - KAFKA_CFG_ZOOKEEPER_CONNECT=zookeeper:2181
    - ALLOW_PLAINTEXT_LISTENER=yes
    - KAFKA_CFG_LISTENER_SECURITY_PROTOCOL_MAP=CLIENT:PLAINTEXT,EXTERNAL:PLAINTEXT
    - KAFKA_CFG_LISTENERS=CLIENT://:9092,EXTERNAL://:29092
    - KAFKA_CFG_ADVERTISED_LISTENERS=CLIENT://kafka:9092,EXTERNAL://localhost:29092
    - KAFKA_CFG_INTER_BROKER_LISTENER_NAME=CLIENT
  depends_on:
    - zookeeper
```

- Klijenti unutar Docker mreže:  
*listener: CLIENT*  
*port: 9092*  
*host: kafka*
- Klijenti van Docker mreže:  
*listener: EXTERNAL*  
*port: 29092*  
*host: localhost*

# Cassandra

```
cassandra:
  image: cassandra:latest
  container_name: cassandra-node
  ports:
    - "9042:9042"
  volumes:
    - cassandra_data:/var/lib/cassandra
```


```
cassandra-setup:
  build: ./cassandra
  image: cassandra-setup
  container_name: cassandra-setup
  links:
    - cassandra
  command: bash -c "sleep 30 && echo Creating Keyspace && cqlsh cassandra -f schema.cql && sleep infinity"
  restart: unless-stopped
```

schema.cql

```
CREATE KEYSPACE IF NOT EXISTS bigdata
WITH REPLICATION = {
  'class' : 'SimpleStrategy',
  'replication_factor' : 1
};

DROP TABLE IF EXISTS bigdata.tripduration;

CREATE TABLE bigdata.tripduration (
  date TIMESTAMP PRIMARY KEY,
  max float,
  min float,
  avg float,
  stddev double,
  taxi1 int,
  count1 int,
  taxi2 int,
  count2 int,
  taxi3 int,
  count3 int
);
```





# Kafka producer

- Python aplikacija
- Lokalno izvršenje
- Čitanje iz .csv datoteke i slanje na Kafka topic

```
Produces message on topic taxiporto:{"trip_id": 1373273114620000046, "call_type": "B", "origin_call": -1, "origin_stand": 47, "taxi_id": 20000046, "start_time": 1373273114, "day_type": "A", "missing_data": "false", "start_lon": -8.654715, "start_lat": 41.173569, "end_lon": -8.651844, "end_lat": 41.187924, "trip_duration": 450.0, "end_time": 1373273564.0}
```

# Apache Spark

- Structured streaming - *micro-batch* obrada tokova podataka  
Izvor podataka: **Kafka**

```
df = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", kafka_url) \
    .option("subscribe", topic) \
    .load()
```

Podaci pročitani sa Kafka topic-a se obrađuju u određenom vremenskom prozoru (1h)

```
grouped_data = df1.groupBy(window("start_time", window_duration))
```

# Apache Spark

Broj taksi vožnji, maksimalna, minimalna, srednja vrednost i standardna devijacija dužine vožnje:

```
statistics = grouped_data.agg(count("*").alias("count"),  
                               avg("trip_duration").alias("avg"),  
                               max("trip_duration").alias("max"),  
                               min("trip_duration").alias("min"),  
                               stddev("trip_duration").alias("stddev"),  
                               col("window.start").alias("start_time"),  
                               col("window.end").alias("end_time")  
                               ).drop("window")
```

# Apache Spark

N najpopularnijih destinacija:

```
grouped_data_by_end_stand = df1.groupby('end_lon', 'end_lat', window('start_time', window_duration))
```

```
popularity = grouped_data_by_end_stand.agg(count("*").alias("count")) \
    .orderBy(desc("count")) \
    .select(col("count"), col("end_lon"), col("end_lat"),
            col("window.start").alias("start_time"), col("window.end").alias("end_time"))

top_taxi_stands = popularity.limit(N)
```

# Apache Spark

Rezultat analize se upisuje u Cassandra bazu

```
query_statistics = statistics.writeStream \  
    .outputMode("update") \  
    .foreachBatch(writeToCassandra) \  
    .start()
```

```
def writeToCassandra(writeDF, epochId):  
    writeDF.write \  
        .format("org.apache.spark.sql.cassandra") \  
        .mode("append") \  
        .options(table="tripduration", keyspace=keyspace) \  
        .save()
```

```
query_popularity = top_taxi_stands.writeStream \  
    .outputMode("complete") \  
    .foreachBatch(writePopularStandsToCassandra) \  
    .start()
```

```
def writePopularStandsToCassandra(writeDF, epochId):  
    writeDF.write \  
        .format("org.apache.spark.sql.cassandra") \  
        .mode("append") \  
        .options(table="popular_stands", keyspace=keyspace) \  
        .save()
```

# Apache Spark

Rezultat analize se upisuje u Cassandra bazu

```
cqlsh:spark_keyspace> SELECT * from popular_stands;
```

start_time	end_time	count	end_lat	end_lon
2013-07-04 16:08:15.000000+0000	2013-07-04 16:08:30.000000+0000	1	41.17057	-8.64813
2013-07-02 15:00:00.000000+0000	2013-07-02 16:00:00.000000+0000	5	41.14983	-8.61961
2013-07-01 02:52:00.000000+0000	2013-07-01 02:53:00.000000+0000	1	41.14239	-8.60587

start_time	end_time	avg	count	max	min	stddev
2013-07-04 11:00:00.000000+0000	2013-07-04 11:00:15.000000+0000	625	3	900	420	247.53787
2013-07-04 11:00:00.000000+0000	2013-07-04 12:00:00.000000+0000	709.96515	287	3315	15	470.53625
2013-07-04 10:11:00.000000+0000	2013-07-04 10:11:15.000000+0000	337.5	2	405	270	95.45942



# Apache Spark

## Spark Master

Status: ALIVE

Workers (2)

Worker Id	Address	State	Cores	Memory	Resources
worker-20230306164844-172.18.0.10-43217	172.18.0.10:43217	ALIVE	8 (8 Used)	14.5 GiB (1024.0 MiB Used)	
worker-20230306164844-172.18.0.9-36319	172.18.0.9:36319	ALIVE	8 (8 Used)	14.5 GiB (1024.0 MiB Used)	


Running Applications (1)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20230307214111-0070	(kill) Taxi Porto 2	16	1024.0 MiB		2023/03/07 21:41:11	root	RUNNING	6.2 min

## Spark Jobs (?)

User: root  
Total Uptime: 8.2 min  
Scheduling Mode: FIFO  
Active Jobs: 2  
Completed Jobs: 193

Event Timeline  
Enable zooming



## App UI

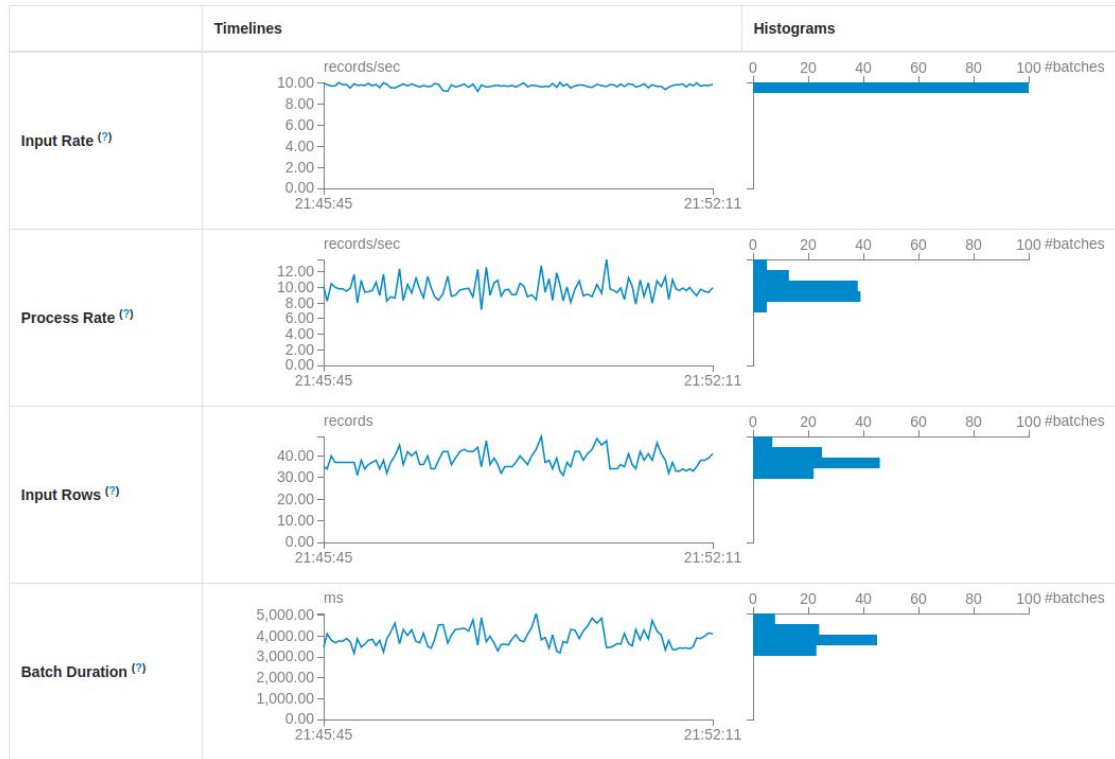
## Streaming Query Statistics

Running batches for **10 minutes 40 seconds** since **2023/03/07 21:41:38** (**143** completed batches)

**Name:** <no name>

**Id:** d361432f-5460-48b6-a9aa-826687f9f35e

**RunId:** 169175a8-68ad-48ef-a5a8-e58b2f947a0e



# Apache Spark

# Apache Flink

- Java Maven aplikacija
- Izvor podataka: Kafka

```
final DeserializationSchema<InputMessage> schema = new DeserializationSchemaInputMessage();

StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();

KafkaSource<InputMessage> source = KafkaSource.<InputMessage>builder()
    .setBootstrapServers(bootstrapServers: "kafka:9092")
    .setTopics(...topics: "taxiporto")
    .setStartingOffsets(OffsetsInitializer.latest())
    .setDeserializer(KafkaRecordDeserializationSchema.valueOnly(schema))
    .build();
```

## Obrada podataka

```
DataStream<InputMessage> ds = env.fromSource(source, WatermarkStrategy.noWatermarks(), sourceName: "Kafka Source").filter(new FilterFunction<InputMessage>() {
    @Override
    public boolean filter(InputMessage value) throws Exception {
        return (value.start_time > 1372636800 && value.start_time < 1375315199);
    }
});
DataStream<TripDurationStatistics> res = ds.windowAll(TumblingProcessingTimeWindows.of(Time.seconds(seconds: 5))).process(new StatisticsProcessWindowFunction());
```

```
CassandraSink.addSink(res)
    .setMapperOptions(() -> new Mapper.Option[] {
        Mapper.Option.saveNullFields(true)
    })
    .setClusterBuilder(new ClusterBuilder() {
        private static final long serialVersionUID = 1L;

        @Override
        protected Cluster buildCluster(Builder builder) {
            return builder.addContactPoints("cassandra-node").withPort(9042).build();
        }
    })
    .build();
env.execute(jobName: "Taxi");
```

Kreiranje Cassandra sink-a

```

public class StatisticsProcessWindowFunction extends ProcessAllWindowFunction<InputMessage, TripDurationStatistics, TimeWindow> {
    @Override
    public void process(ProcessAllWindowFunction<InputMessage, TripDurationStatistics, TimeWindow>.Context context,
        Iterable<InputMessage> elements, Collector <TripDurationStatistics> out) throws Exception {
        float sum = 0;
        float max = 0;
        float min = 50000;
        float avg = 0;
        double stddev = 0;
        float count = 0;

        HashMap<Integer, Integer> taxis = new HashMap<>();

        for (InputMessage msg : elements) {
            count++;
            sum += msg.trip_duration;
            if (msg.trip_duration > max)
                max = msg.trip_duration;
            if (msg.trip_duration < min)
                min = msg.trip_duration;
            if (!taxis.containsKey(msg.taxi_id)) {
                taxis.put(msg.taxi_id, 1);
            }
            else {
                int val = taxis.get(msg.taxi_id) + 1;
                taxis.replace(msg.taxi_id, val);
            }
        }

        int taxi1 = (int)taxis.keySet().toArray()[0];
        int count1 = taxis.get(taxi1);
        int taxi2 = (int)taxis.keySet().toArray()[1];
        int count2 = taxis.get(taxi2);
        int taxi3 = (int)taxis.keySet().toArray()[2];
        int count3 = taxis.get(taxi3);

        avg = sum / count;
        for (InputMessage msg : elements) {
            stddev += Math.pow(msg.trip_duration - avg, 2);
        }

        stddev = Math.sqrt(stddev / count);
        Date date = new Date();

        TripDurationStatistics res = new TripDurationStatistics(min, max, avg, stddev, date, taxi1, count1, taxi2, count2, taxi3, count3);
        out.collect(res);
    }
}

```

Klasa i funkcija za  
analizu toka  
podataka

# Apache Flink

*Cassandra* - prikaz podataka

```
cqlsh:bigdata> SELECT * from tripduration;
```

date	avg	count1	count2	count3	max	min	stddev	taxi1	taxi2	taxi3
2023-03-07 23:57:55.001000+0000	594.375	2	2	1	1365	15	322.81416	20000688	20000432	20000307
2023-03-07 23:56:00.001000+0000	937.65625	1	1	1	6180	7.5	918.04244	20000177	20000113	20000050
2023-03-07 23:56:10.000000+0000	797.8125	1	1	1	2205	45	489.22441	20000113	20000177	20000688
2023-03-07 23:55:10.001000+0000	680.35712	1	1	2	3300	7.5	616.98657	20000561	20000560	20000496



# Apache Flink


1 Job Manager  
2 Task Manager-a


Job submit

Uploaded Jars

+ Add New

Name	Upload Time	Entry Class	
BigData.jar	2023-03-08, 00:26:28	com.flink.App	<a href="#">Delete</a>

 com.flink.App

 Program Arguments

☐ Allow Non Restored State

↺ 2

📁 Savepoint Path

Show Plan

Submit

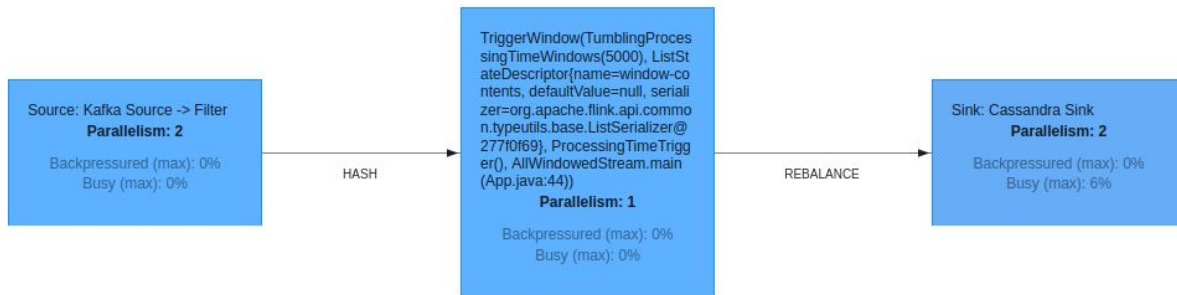
## Task Managers

Path, ID	Data Port	Last Heartbeat	All Slots	Free Slots	CPU Cores	Physical MEM	JVM Heap Size	Flink Managed MEM
<a href="#">172.18.0.16:42515-aa6d05</a> akka.tcp://flink@172.18.0.16:42515/user/rpc/taskmanager_0	32937	2023-03-08 00:53:45	1	0	8	15.5 GB	512 MB	512 MB
<a href="#">172.18.0.17:32933-9998d0</a> akka.tcp://flink@172.18.0.17:32933/user/rpc/taskmanager_0	43667	2023-03-08 00:53:45	1	0	8	15.5 GB	512 MB	512 MB

# Apache Flink

## Running Jobs

Job Name	Start Time	Duration	End Time	Tasks	Status
Taxi	2023-03-08 00:50:02	4m 12s	—	5 / 5	RUNNING



Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Start Time	Duration	Tasks
Source: Kafka Source -> Filter	RUNNING	0 B	0	0 B	199	2	2023-03-08 00:50:04	1m 1s	2
TriggerWindow(TumblingProcessingTimeWindows(5000), ListStateD...	RUNNING	18.2 KB	198	0 B	3	1	2023-03-08 00:50:04	1m 1s	1
Sink: Cassandra Sink	RUNNING	239 B	3	0 B	0	2	2023-03-08 00:50:04	1m 1s	2

# Projekat 3

- Spark aplikacija za treniranje modela mašinskog učenja
  - Spark Streaming aplikacija
  - InfluxDB
  - Grafana
-

# Spark ML

- Predikcija trajanja taksi vožnje
- *Gradient-Boosted Tree* (`pyspark.ml.regression.GBTRegressor`)

- Skaliranje podataka korišćenjem robust scaler-a

```
split = va.randomSplit([0.8, 0.2], 3333)
scaler = RobustScaler().setInputCol('features').setOutputCol('scaled_features')
scaler_model = scaler.fit(split[0])
scaler_model.write().overwrite().save(SCALER_LOCATION)
train = scaler_model.transform(split[0])
test = scaler_model.transform(split[1])
```

- Optimizacija hiperparametara (`pyspark.ml.tuning.ParamGridBuilder`)

```
paramGrid = ParamGridBuilder()\
    .addGrid(gbt.maxBins, [100, 200])\
    .addGrid(gbt.maxDepth, [2, 4, 10])\
    .build()
```

# Spark ML

```
test = scaler_model.transform(split[1])

gbt = GBRegressor(featuresCol='scaled_features', labelCol='trip_duration', maxIter=10)

pipeline = Pipeline(stages=[gbt])

paramGrid = ParamGridBuilder()\
    .addGrid(gbt.maxBins, [100, 200])\
    .addGrid(gbt.maxDepth, [2, 4, 10])\
    .build()

evaluator = RegressionEvaluator(labelCol='trip_duration', predictionCol='prediction', metricName='rmse')

cv = CrossValidator().setEstimator(pipeline).setEvaluator(evaluator).setEstimatorParamMaps(paramGrid).setNumFolds(5)

mp = cv.fit(train)

prediction = mp.transform(test)
prediction.show()
```

# Spark ML

- Evaluacija

```
evaluator = RegressionEvaluator(labelCol='trip_duration', predictionCol='prediction', metricName='r2')
r2 = evaluator.evaluate(prediction)
evaluator = RegressionEvaluator(labelCol='trip_duration', predictionCol='prediction', metricName='mae')
mae = evaluator.evaluate(prediction)
```

- Model se čuva na HDFS

## Browse Directory

/data/model

Go!

Show

25

entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Mar 02 23:14	0	0 B	<a href="#">bestModel</a>	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Mar 02 23:14	0	0 B	<a href="#">estimator</a>	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Mar 02 23:14	0	0 B	<a href="#">evaluator</a>	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Mar 02 23:14	0	0 B	<a href="#">metadata</a>	

# Spark Streaming aplikacija

- Izvor podataka: Kafka topic
- Model se učitava sa HDFS-a

```
appName = "Taxi Porto"
spark = SparkSession.builder.appName(appName).getOrCreate()
spark.sparkContext.setLogLevel("ERROR")
df = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", KAFKA_URL) \
    .option("subscribe", "taxiporto") \
    .load()
df.printSchema()

df1 = df.selectExpr("CAST(value AS STRING)").select(F.from_json(F.col("value"), schema).alias("data")).select("data.*")

for col_name in nullable_columns:
    df1 = df1.withColumn(col_name, F.when(F.isnull(df1[col_name]), 0).otherwise(df1[col_name]).alias(col_name))
df1 = df1.drop('missing_data')
df1 = df1.drop('trip_id')
df1 = df1.drop('day_type')

indexed = indexer.transform(df1)

va = VectorAssembler().setInputCols(columns).setOutputCol('features').setHandleInvalid("skip").transform(indexed)

scaled = scaler.transform(va)

prediction = model.transform(scaled)

prediction.printSchema()

query = prediction.writeStream \
    .foreach(InfluxDBWriter()) \
    .start()
query.awaitTermination()
```

```
model = CrossValidatorModel.load(MODEL_LOCATION)
scaler = RobustScalerModel.load(SCALER_LOCATION)
indexer = StringIndexerModel.load(INDEXER_LOCATION)
```



# InfluxDB

```
class InfluxDBWriter:
    def __init__(self):
        self._org = 'taxiportodb'
        self._token = '2c83186a-caab-425a-9594-9d4c00544939'
        self.client = InfluxDBClient(
            url = "http://influxdb:8086", token=self._token, org = self._org)
        self.write_api = self.client.write_api(write_options=SYNCHRONOUS)

    def open(self, partition_id, epoch_id):
        print("Opened %d, %d" % (partition_id, epoch_id))
        return True


    def process(self, row):
        self.write_api.write(bucket='taxiportodb',
                             record=self._row_to_line_protocol(row))

    def close(self, error):
        self.write_api._del__()
        self.client._del__()
        print("Closed with error: %s" % str(error))

    def _row_to_line_protocol(self, row):
        print(row)
        timestamp = datetime.utcnow().strftime('%Y-%m-%dT%H:%M:%SZ')
        return Point.measurement(KAFKA_TOPIC).tag("measure", KAFKA_TOPIC) \
            .field("start_lat", float(row['start_lat'])) \
            .field("end_lat", float(row['start_lat'])) \
            .field("start_lon", float(row['start_lon'])) \
            .field("end_lon", float(row['end_lon'])) \
            .field("trip_duration", float(row['trip_duration'])) \
            .field("prediction", float(row['prediction'])) \
            .time(timestamp, write_precision='ms')
```

```
influxdb:
  image: influxdb:1.8
  ports:
    - "0.0.0.0:8086:8086"
  container_name: influxdb
  hostname: influxdb
  environment:
    - INFLUXDB_ADMIN_ENABLED=true
    - INFLUXDB_DB=taxiportodb
    - DOCKER_INFLUXDB_INIT_ADMIN_TOKEN=2c83186a-caab-425a-9594-9d4c00544939
```

# Vizuelizacija u Grafani



# Vizualizacija u Grafani

